

Importing Libraries

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import os
```

Importing Deep Learning Libraries


```
from tensorflow.keras.preprocessing.image import load_img,
img_to_array
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import
Dense, Input, Dropout, GlobalAveragePooling2D, Flatten, Conv2D, BatchNormali
zation, Activation, MaxPooling2D
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.optimizers import Adam, SGD, RMSprop
```

```
2021-09-12 10:12:57.352482: I
tensorflow/stream_executor/platform/default/dso_loader.cc:49]
Successfully opened dynamic library libcudart.so.11.0
```

Displaying Images


```
picture_size = 48
folder_path = "../input/face-expression-recognition-dataset/images/"
expression = 'happy'

plt.figure(figsize= (12,12))
for i in range(1, 10, 1):
    plt.subplot(3,3,i)
    img = load_img(folder_path+"train/"+expression+"/"+
                    os.listdir(folder_path + "train/" + expression)[i],
                    target_size=(picture_size, picture_size))
    plt.imshow(img)
plt.show()
```



```
expression = 'disgust'

plt.figure(figsize= (12,12))
for i in range(1, 10, 1):
    plt.subplot(3,3,i)
    img = load_img(folder_path+"train/"+expression+"/"+
                    os.listdir(folder_path + "train/" + expression)[i],
                    target_size=(picture_size, picture_size))
    plt.imshow(img)
plt.show()
```



Making Training and Validation Data

```
batch_size=128
datagen_train=ImageDataGenerator()
datagen_val=ImageDataGenerator()

train_set= datagen_train.flow_from_directory(folder_path+"train",
target_size=(picture_size,picture_size),
color_mode="grayscale",
batch_size=batch_size,
class_mode="categorical",
```

```

                                shuffle=True)
test_set=datagen_val.flow_from_directory(folder_path+"validation",
target_size=(picture_size,picture_size),
                                color_mode="grayscale",
                                batch_size=batch_size,
                                class_mode="categorical",
                                shuffle=False)

```

Found 28821 images belonging to 7 classes.
Found 7066 images belonging to 7 classes.

Model Building

```

no_of_classes=7

model=Sequential()

#1st CNN layer
model.add(Conv2D(64,(3,3),padding="same",input_shape=(48,48,1)))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

#2nd CNN layer
model.add(Conv2D(128,(5,5),padding="same"))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

#3rd CNN layer
model.add(Conv2D(512,(3,3),padding="same"))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

#4th CNN layer
model.add(Conv2D(512,(3,3), padding='same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())

```

```

#Fully connected 1st layer
model.add(Dense(256))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(Dropout(0.25))

# Fully connected layer 2nd layer
model.add(Dense(512))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(Dropout(0.25))

model.add(Dense(no_of_classes,activation="softmax"))

opt = Adam(lr = 0.0001)
model.compile(optimizer=opt,loss='categorical_crossentropy',
metrics=['accuracy'])
model.summary()

```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_7 (Conv2D)	(None, 48, 48, 64)	640
batch_normalization_9 (Batch Normalization)	(None, 48, 48, 64)	256
activation_9 (Activation)	(None, 48, 48, 64)	0
max_pooling2d_7 (MaxPooling2D)	(None, 24, 24, 64)	0
dropout_9 (Dropout)	(None, 24, 24, 64)	0
conv2d_8 (Conv2D)	(None, 24, 24, 128)	204928
batch_normalization_10 (Batch Normalization)	(None, 24, 24, 128)	512
activation_10 (Activation)	(None, 24, 24, 128)	0
max_pooling2d_8 (MaxPooling2D)	(None, 12, 12, 128)	0
dropout_10 (Dropout)	(None, 12, 12, 128)	0
conv2d_9 (Conv2D)	(None, 12, 12, 512)	590336
batch_normalization_11 (Batch Normalization)	(None, 12, 12, 512)	2048
activation_11 (Activation)	(None, 12, 12, 512)	0
max_pooling2d_9 (MaxPooling2D)	(None, 6, 6, 512)	0

dropout_11 (Dropout)	(None, 6, 6, 512)	0
conv2d_10 (Conv2D)	(None, 6, 6, 512)	2359808
batch_normalization_12 (Batch Normalization)	(None, 6, 6, 512)	2048
activation_12 (Activation)	(None, 6, 6, 512)	0
max_pooling2d_10 (MaxPooling2D)	(None, 3, 3, 512)	0
dropout_12 (Dropout)	(None, 3, 3, 512)	0
flatten_1 (Flatten)	(None, 4608)	0
dense_2 (Dense)	(None, 256)	1179904
batch_normalization_13 (Batch Normalization)	(None, 256)	1024
activation_13 (Activation)	(None, 256)	0
dropout_13 (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 512)	131584
batch_normalization_14 (Batch Normalization)	(None, 512)	2048
activation_14 (Activation)	(None, 512)	0
dropout_14 (Dropout)	(None, 512)	0
dense_4 (Dense)	(None, 7)	3591
=====		
Total params: 4,478,727		
Trainable params: 4,474,759		
Non-trainable params: 3,968		

Fitting the Model with Training and Validation Data

```
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau

checkpoint = ModelCheckpoint("./model.h5", monitor='val_acc',
                             verbose=1, save_best_only=True, mode='max')
```

```

early_stopping = EarlyStopping(monitor='val_loss',
                                min_delta=0,
                                patience=3,
                                verbose=1,
                                restore_best_weights=True
                                )

reduce_learningrate = ReduceLRonPlateau(monitor='val_loss',
                                          factor=0.2,
                                          patience=3,
                                          verbose=1,
                                          min_delta=0.0001)

callbacks_list = [early_stopping,checkpoint,reduce_learningrate]

epochs = 48

model.compile(loss='categorical_crossentropy',
              optimizer = Adam(lr=0.001),
              metrics=['accuracy'])

history = model.fit_generator(generator=train_set,
                              steps_per_epoch=train_set.n//train_set.batch_size,
                              epochs=epochs,
                              validation_data = test_set,
                              validation_steps =
test_set.n//test_set.batch_size,
                              callbacks=callbacks_list
                              )

```

```

/opt/conda/lib/python3.7/site-packages/tensorflow/python/keras/
engine/training.py:1844: UserWarning: `Model.fit_generator` is
deprecated and will be removed in a future version. Please use
`Model.fit`, which supports generators.

```

```

    warnings.warn("`Model.fit_generator` is deprecated and '
2021-09-12 10:29:59.646073: I
tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:116] None of
the MLIR optimization passes are enabled (registered 2)
2021-09-12 10:29:59.650537: I
tensorflow/core/platform/profile_utils/cpu_utils.cc:112] CPU
Frequency: 2199995000 Hz

```

Epoch 1/48

```

2021-09-12 10:30:01.184226: I
tensorflow/stream_executor/platform/default/dso_loader.cc:49]
Successfully opened dynamic library libcublas.so.11
2021-09-12 10:30:02.080952: I
tensorflow/stream_executor/platform/default/dso_loader.cc:49]
Successfully opened dynamic library libcublasLt.so.11

```

```

2021-09-12 10:30:02.118051: I
tensorflow/stream_executor/platform/default/dso_loader.cc:49]
Successfully opened dynamic library libcudnn.so.8

225/225 [=====] - 165s 700ms/step - loss:
1.9170 - accuracy: 0.2713 - val_loss: 1.5695 - val_accuracy: 0.4071
Epoch 2/48
225/225 [=====] - 30s 132ms/step - loss:
1.4542 - accuracy: 0.4382 - val_loss: 1.3962 - val_accuracy: 0.4744
Epoch 3/48
225/225 [=====] - 30s 131ms/step - loss:
1.2821 - accuracy: 0.5069 - val_loss: 1.2841 - val_accuracy: 0.5138
Epoch 4/48
225/225 [=====] - 30s 133ms/step - loss:
1.1848 - accuracy: 0.5479 - val_loss: 1.1735 - val_accuracy: 0.5614
Epoch 5/48
225/225 [=====] - 30s 133ms/step - loss:
1.1237 - accuracy: 0.5754 - val_loss: 1.3103 - val_accuracy: 0.4747
Epoch 6/48
225/225 [=====] - 29s 129ms/step - loss:
1.0747 - accuracy: 0.5898 - val_loss: 1.3306 - val_accuracy: 0.4913
Epoch 7/48
225/225 [=====] - 29s 130ms/step - loss:
1.0090 - accuracy: 0.6181 - val_loss: 1.0748 - val_accuracy: 0.5972
Epoch 8/48
225/225 [=====] - 29s 129ms/step - loss:
0.9811 - accuracy: 0.6287 - val_loss: 1.1196 - val_accuracy: 0.5778
Epoch 9/48
225/225 [=====] - 29s 127ms/step - loss:
0.9303 - accuracy: 0.6496 - val_loss: 1.0919 - val_accuracy: 0.5876
Epoch 10/48
225/225 [=====] - 29s 128ms/step - loss:
0.8880 - accuracy: 0.6676 - val_loss: 1.1163 - val_accuracy: 0.5786
Restoring model weights from the end of the best epoch.

Epoch 00010: ReduceLRonPlateau reducing learning rate to
0.000200000000949949026.
Epoch 00010: early stopping

model.save_weights('face_emotion_model.h5')

```

Plotting Accuracy & Loss

```

plt.style.use('dark_background')

plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : Adam', fontsize=10)

```



```

plt.ylabel('Loss', fontsize=16)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()

```

